


# **USING ATTENTION FOR SPEECH RECOGNITION AND NLP**

# Attention ASR at a Glance




“Language model”  
RNN Generates text  
letter-by-letter

Alignment model:  
Attention mechanism

“Acoustic model”  
Convolutional and  
recurrent layers

Speech features  
Mel spectrogram

# Attention ASR at a Glance



“Language model”  
RNN Generates text  
letter-by-letter

Alignment model:  
Attention mechanism


“Acoustic model”  
Convolutional and  
recurrent layers

Speech features  
Mel spectrogram

# Attention ASR at a Glance


Network defines  
 $p(\text{Words}|\text{Audio}; \Theta)$   
where  
 $\Theta$  are parameters.

Training uses  
gradient  
optimization



“Language model”  
RNN Generates text  
letter-by-letter



Alignment model:  
Attention mechanism



“Acoustic model”  
Convolutional and  
recurrent layers


Speech features  
Mel spectrogram

# Attention Mechanism in Action




# Challenges

- Overconfidence.
- Long sequences and repetitions.
- Language model integration and coverage.




# Overconfidence

Ground truth, total log probability -25



Beam search result: total log probability -3.7



# Key Observations


- Accurate next-step predictions:  
99.9% train/96% test
- Overconfidence:  
 $p(\text{first guess}) \gg p(\text{second guess})$
- A “second guess” of the net costs as much as several “first guess” predictions
  - Beam search ineffective at large beams
  - Very hard to balance decoding costs (e.g. LM)



# A Simple Experiment

- After training, tweak SoftMax temperature

$$\text{SoftMax}(Y) = \frac{\exp(Y_i/T)}{\sum_j \exp(Y_j/T)}$$



# Training With 1-hot Labels

- The cross-entropy cost for one utterance

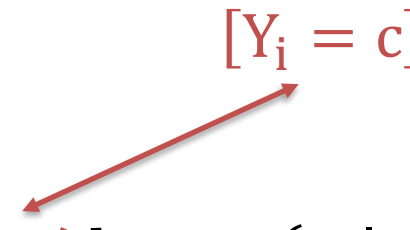
$$-\sum_{i=1}^N \sum_c [\mathbf{Y}_i = \mathbf{c}] \log p_{\Theta}(Y_i | Y_{<i}, X_i)$$

- When model is 99% accurate...
- The only way to reduce cost is to make  $p_{\Theta}(Y_i | Y_{<i}, X_i)$  a Dirac delta...

# Training With Label Smoothing

- Introduced in Inception V2 (arXiv:1512.00567)

- Change the cost to:

$$-\frac{1}{N} \sum_{i=1}^N \sum_{c=1}^C T(Y_i, c) \log p_{\Theta}(Y_i | X_i)$$


$[Y_i = c]$

- $T(Y_i, c)$  is a smoothing distribution, e.g.

$$T(Y_i, c) = \begin{cases} \beta, & \text{when } Y_i = c \\ \frac{1 - \beta}{C - 1}, & \text{otherwise} \end{cases}$$

- Even better: smooth the  $1 - \beta$  according to class marginal probabilities (unigrams)


# Effects of Label Smoothing

- Reduces overconfidence and regularizes
- Also prevents gradient vanishing:
  - Without smoothing SoftMax derivative is  $p_{\Theta}(Y_i|X_i) - [Y_i = c]$
  - This vanishes when  $p_{\Theta}(Y_i|X_i) \approx 1$
  - Effectively the model stops training on correctly classified characters

# Label Smoothing vs Other Regularizers

At a high level, all regularizers want to forbid large changes of output for small changes of input.

- E.g. weight decay




Magnitude of  $w_o$  controls the output sensitivity  $\frac{\partial o}{\partial L_o} = w_o^T$

- Label smoothing may be easier to use:
  - Easy to say how smooth the output should be
  - Hard to say how large the weights should be


# Effects of Label smoothing

- Regularization (next character accuracy increase 96% -> 97%)
- Increase of neg log-probability of best predictions -> other costs easier to balance



# SoftMax Temperature and Label Smoothing

- Temperature tweaking no longer needed:



# Trouble With Long Sequences

A simple experiment:

1. Train a network as usual.
2. Concatenate test utterances a few times.
3. Decode as usual.

Performance drops dramatically.

On long utterances decoding completely fails.




# Investigation of Long Inputs

The setup:

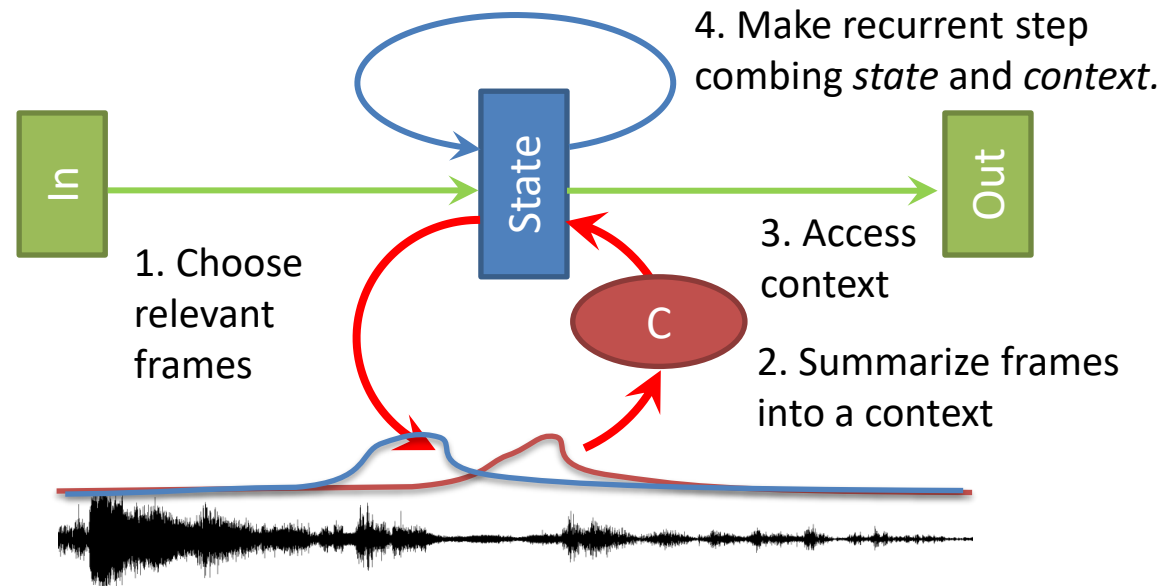
- concatenate utterances
- do force alignment (feed the correct inputs)

Typical result



Our hypothesis: the net learns an implicit location encoder. It is not robust to long utterances.


# Location-aware Attention



- We want to separate repetitions of the same sound
- Use the selection from the last step to make the new selection
- This enables the model to learn concepts like “later than last” or “close to last”.

# Location-aware attention helps

- Decoding error rate increases from 18% to 20%



- One more “trick”: constrain the attention mechanism to select only few frames
  - Keep up to  $K$  with highest scores
  - Limit selection to the vicinity of previous one

# Decoding With Language Models

- Extend the beam search cost

$$\hat{Y} = \arg \min_Y -\log p_{\Theta}(Y|X) - \alpha p_{LM}(Y)$$

Transcript	LM cost $\log p(y)$	Model cost $\log p(y x)$	
"chase is nigeria's registrar and the society is an independent organization hired to count votes"	-108.5	-34.5	Ground truth
"in the society is an independent organization hired to count votes"	-64.6	-19.9	Decoded
"chase is nigeria's registrar"	-40.6	-31.2	Severe Transcript Truncation
"chase's nature is register"	-37.8	-20.3	
""	-3.5	-12.5	

# Promoting long transcripts

Seems easy:

$$\hat{Y} = \arg \min_Y -\log p_{\Theta}(Y|X) - \alpha p_{LM}(Y) - \beta |Y|$$

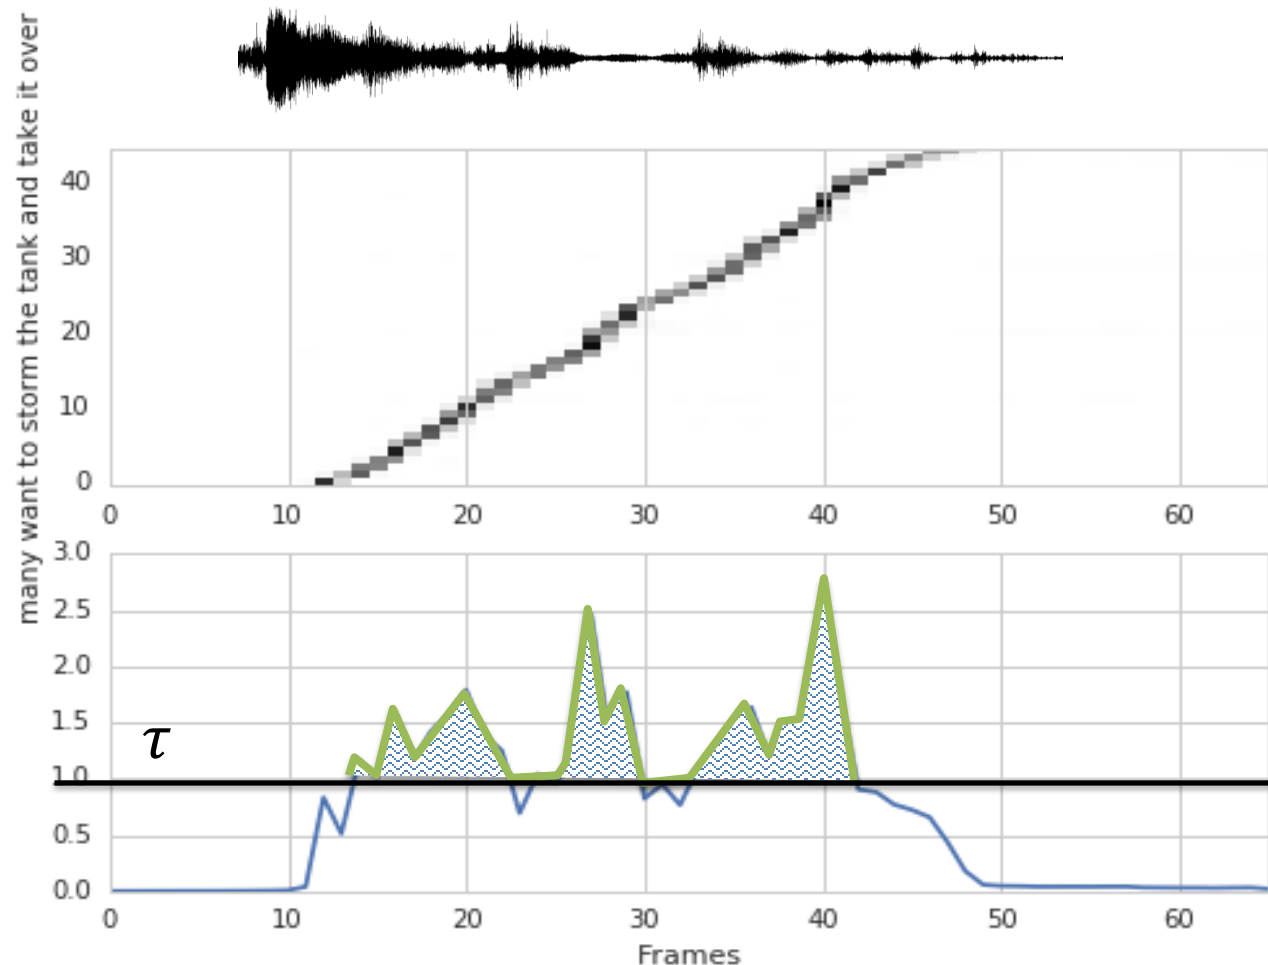
Problem: if any sequence of characters is cheap and the cost becomes negative, the model will keep repeating itself...

# Coverage Criterion

Force decoding of all frames,  
but prevent looping.

$$\text{coverage} = \sum_f [\sum_i \alpha_{fi} > \tau]$$



Can't loop: a frame is counted  
at most once



# **BEYOND SIMPLE SPEECH RECOGNITION**

# Speech-to-text translation

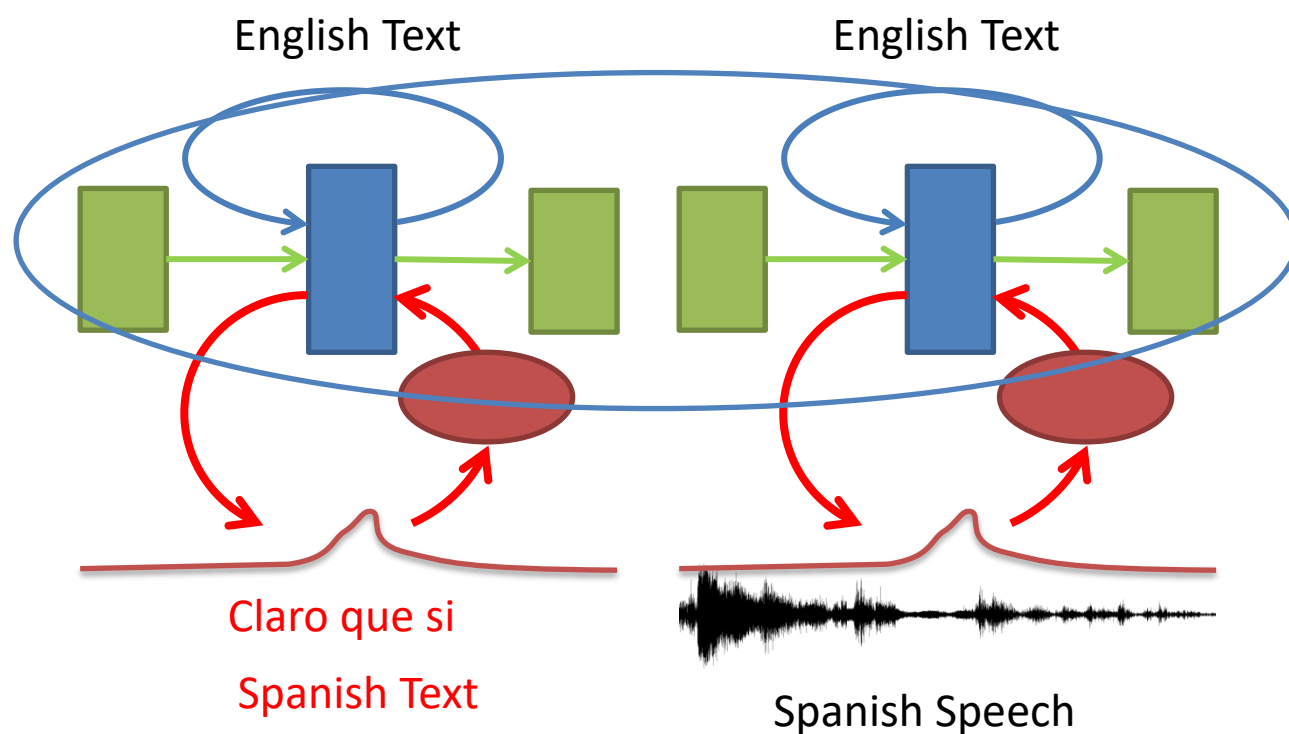
- Seq2seq model





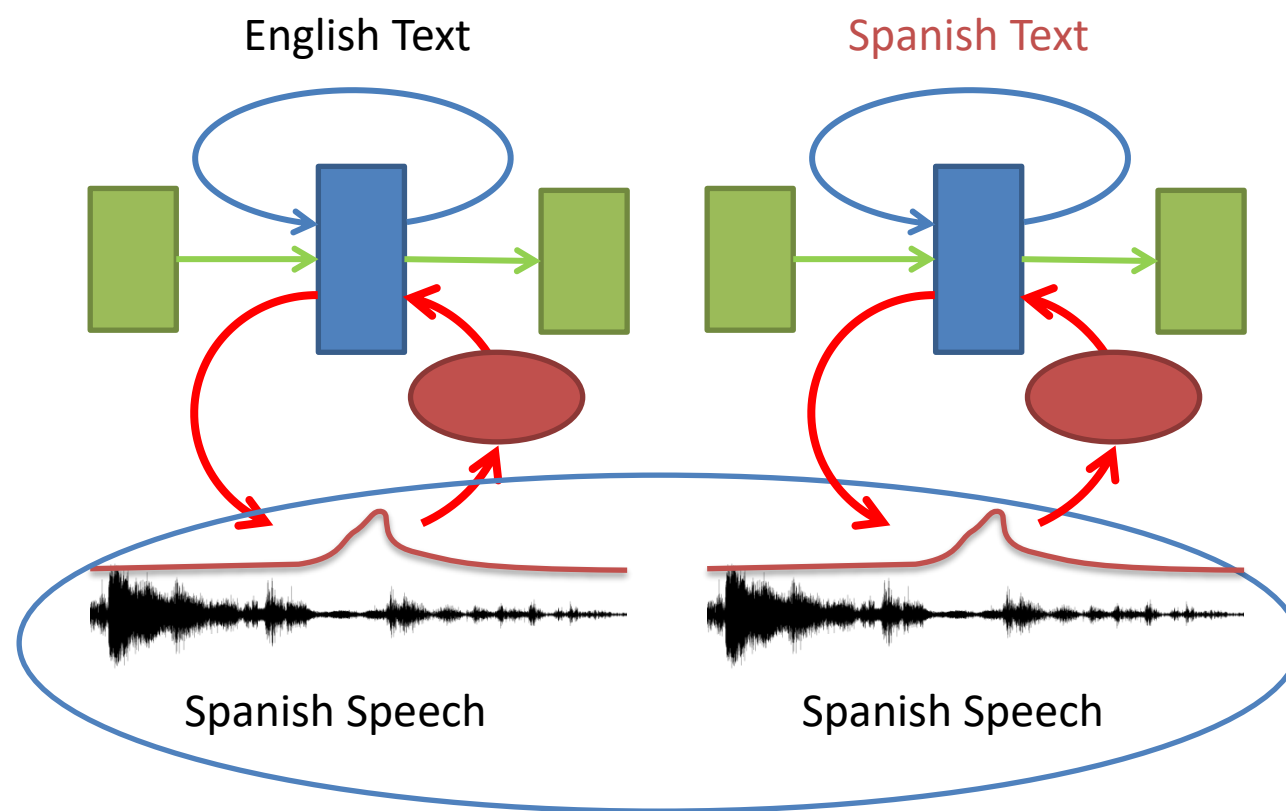
# Multitask Learning, or Exploit All Data

Share weights of the decoder, separate encoders

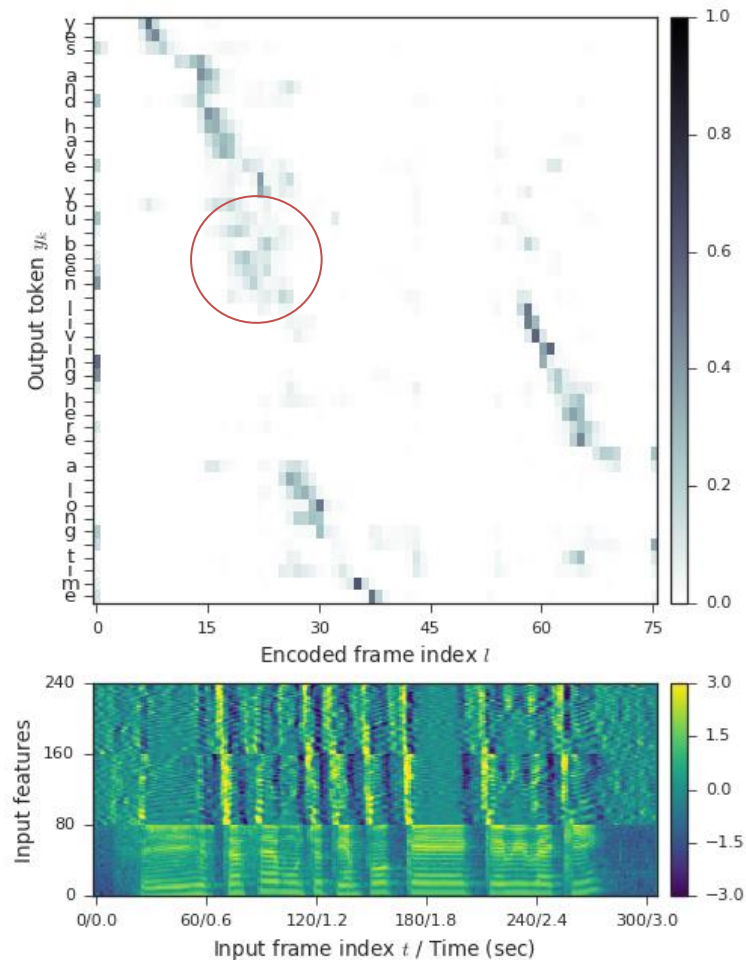
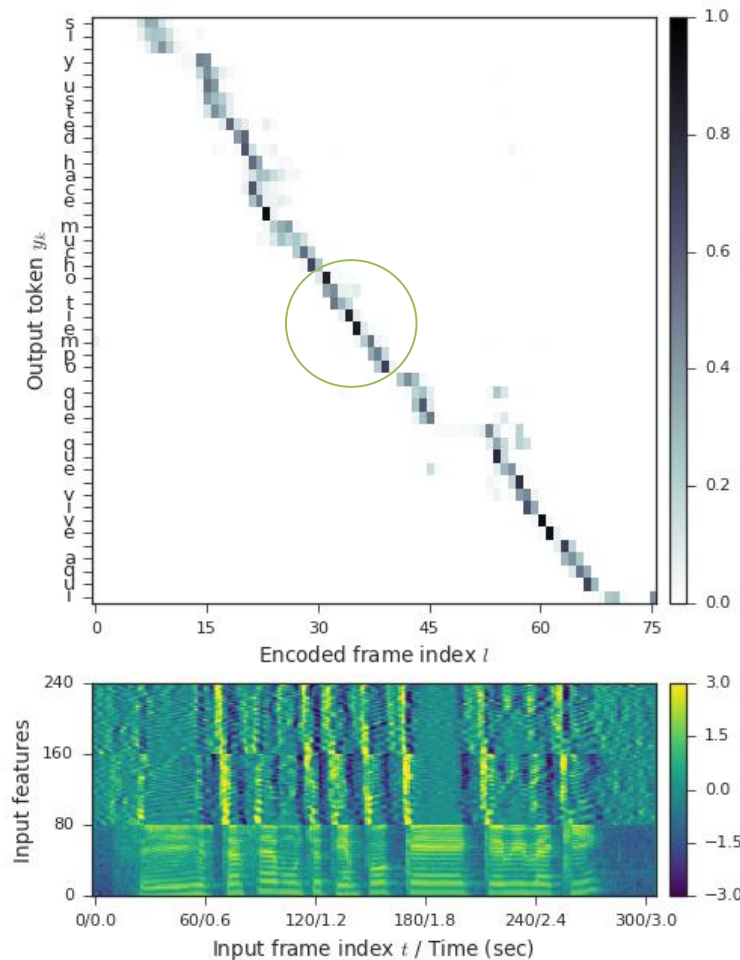


# Multitask Learning, or Exploit All Data

Share weights of the encoder, separate decoders

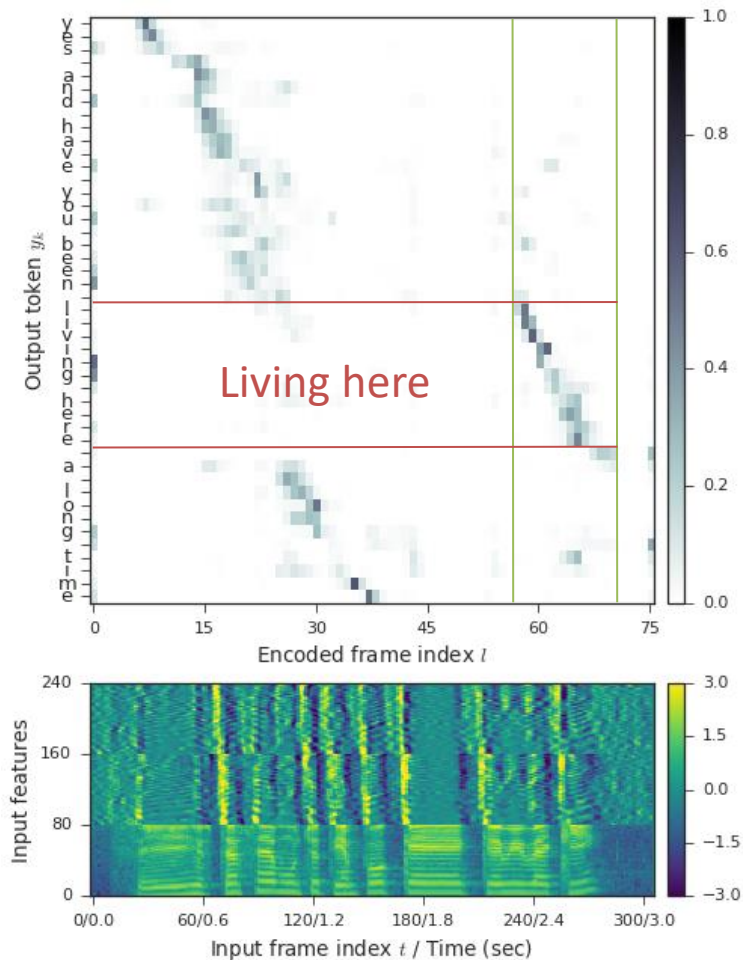
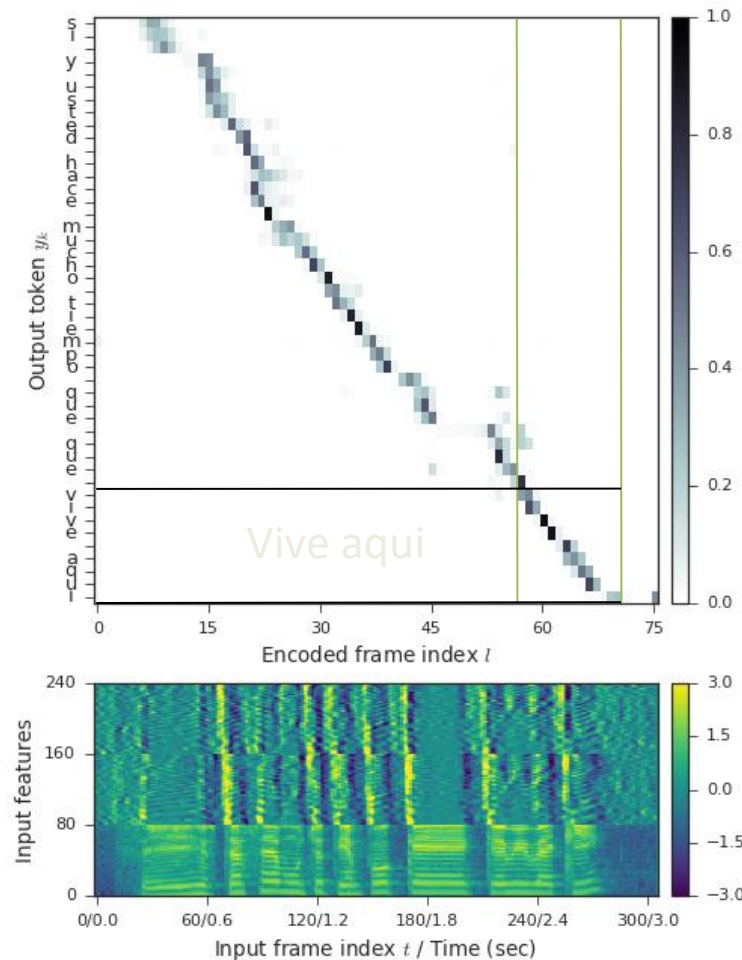


# Seq2seq Speech Translation: Attention



- recognition attention very **confident**
- translation attention **smoothed** out across many spectrogram frames for each output character
  - ambiguous mapping between Spanish speech acoustics and English text

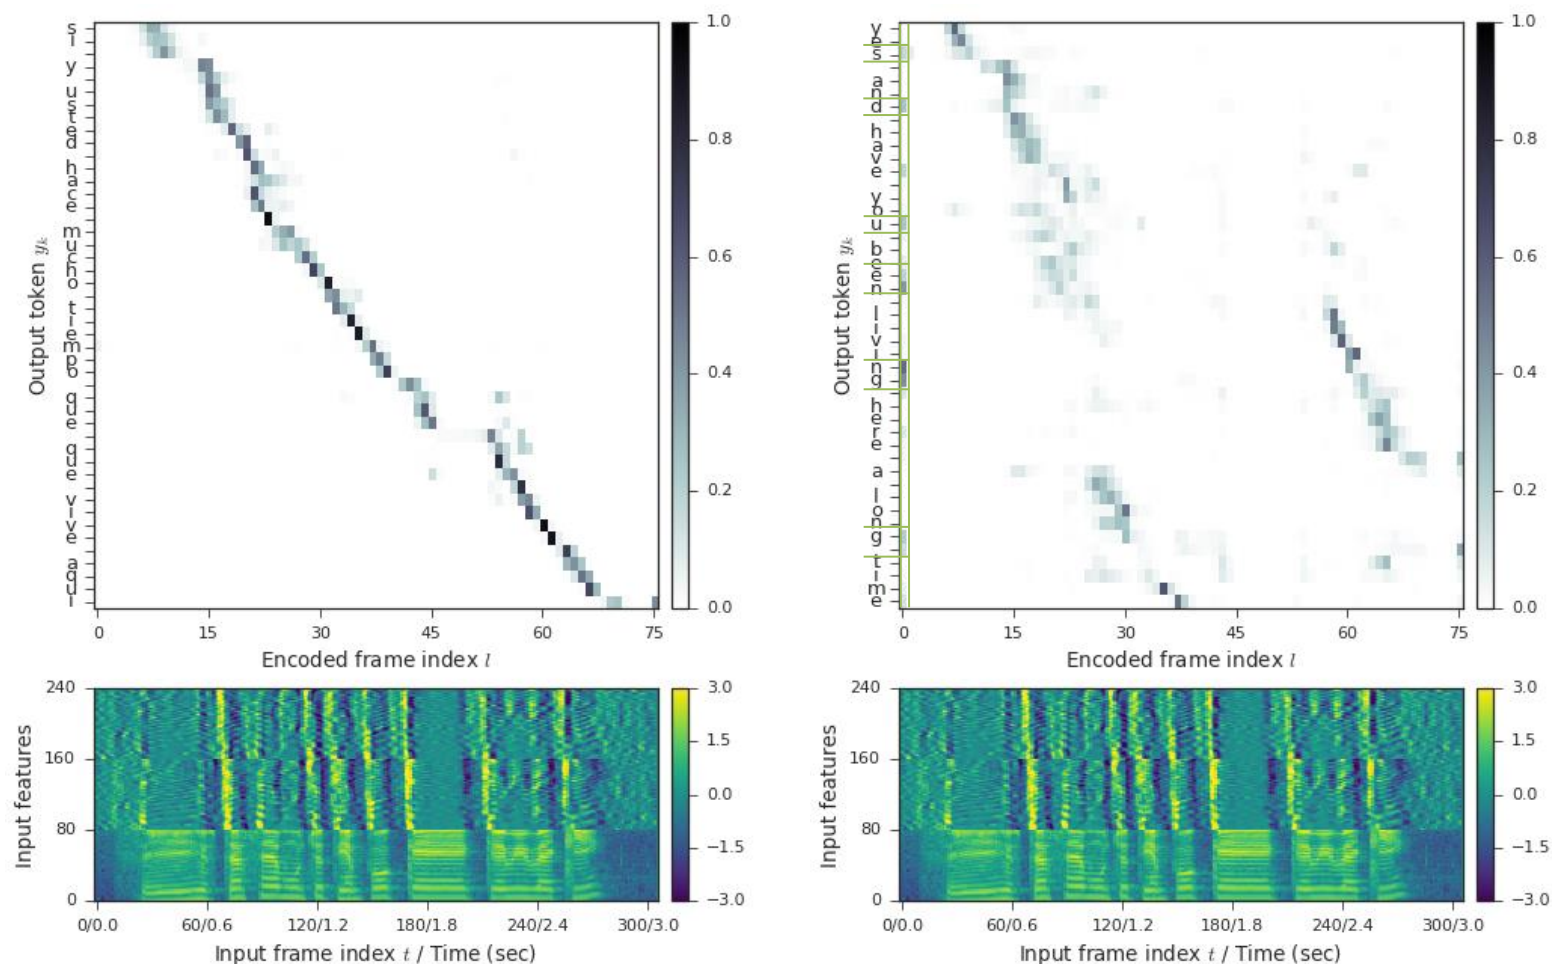
# Seq2seq Speech Translation: Attention



- speech recognition attention is mostly monotonic
- translation attention reorders input: **same frames** attended to for "vive aqui" and "living here"

Weiss, Chorowski et al., Sequence-to-Sequence Models Can Directly Translate Foreign Speech, INTERSPEECH 2017

# Seq2seq Speech Translation: Example attention



translation model **attends to the beginning of input** (i.e. silence) for the last few letters in each word

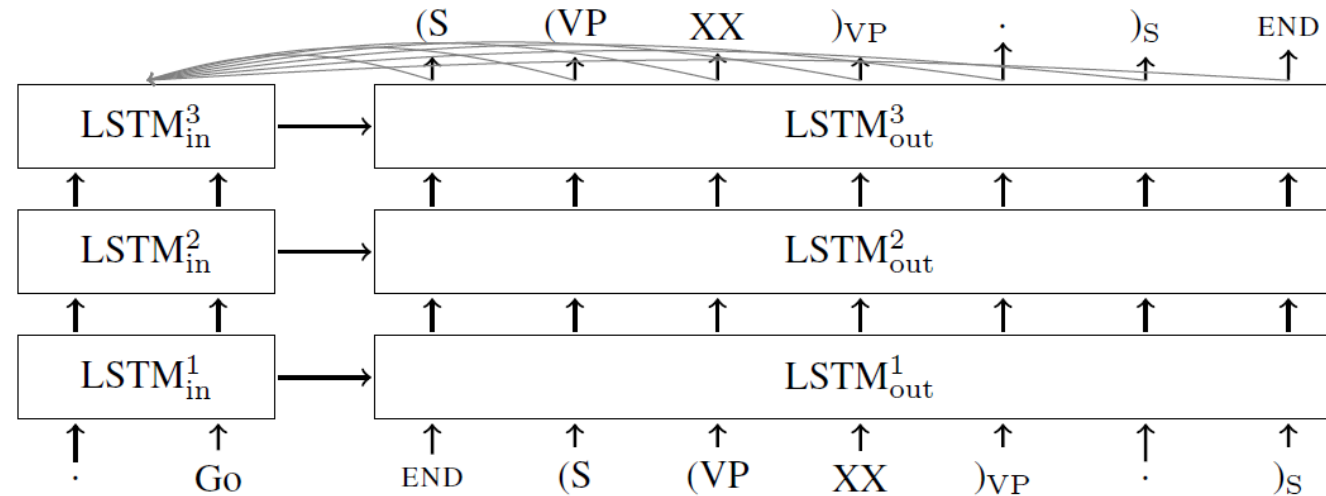

- already made a decision about word to emit, just acts a language model to spell it out.

# End-to-end systems in NLP:

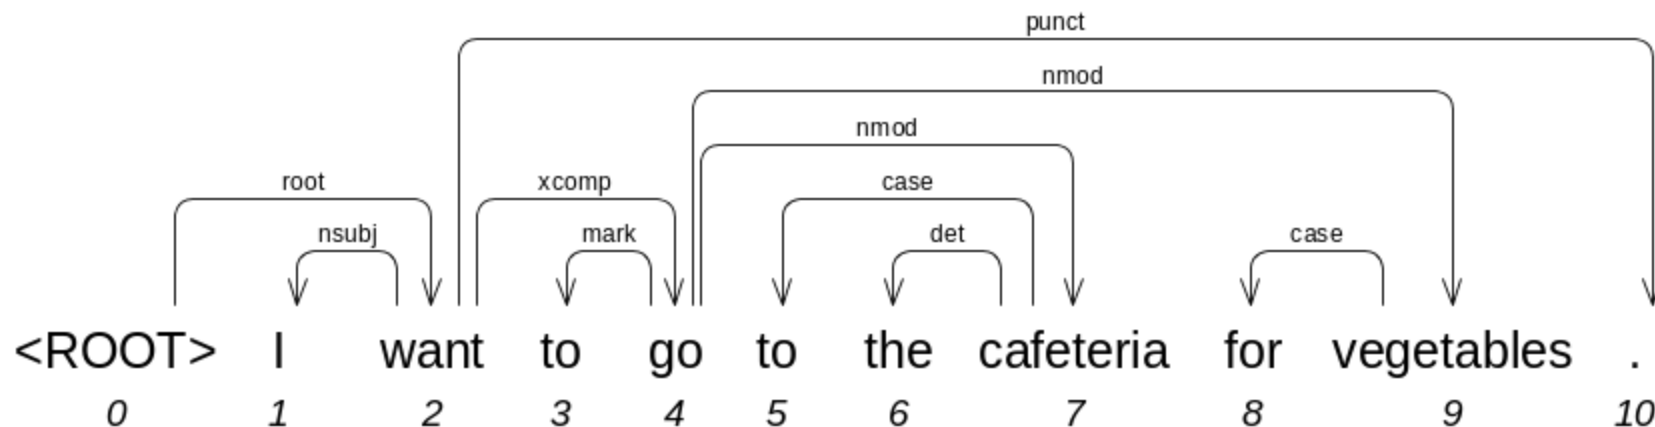
## How to parse sentences?

For constituency parsing:  
Treat parsing as a sequence-to-sequence problem:

- Input: sentence  
„Go .”
- Output: linearized parse tree:  
„(S (VP XX )VP . )S END”



# Dependency parsing

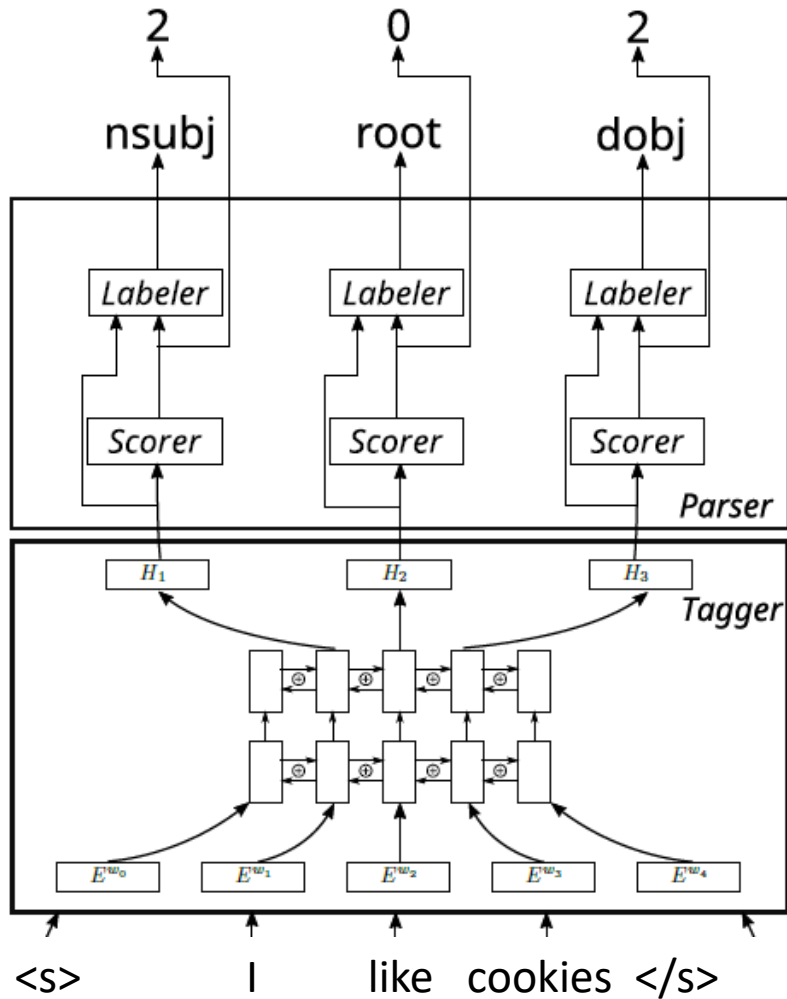


- Desired output: directed edges between words.
- At each step the attention selects a few words.
- Idea: use the selection weights as pointers.

Chorowski et al. "Read, Tag, and Parse All at Once, or Fully-neural Dependency Parsing",  
arxiv <https://arxiv.org/pdf/1609.03441>

Zapotocny et al. "On Multilingual Training of Neural Dependency Parsers" TSD 2017

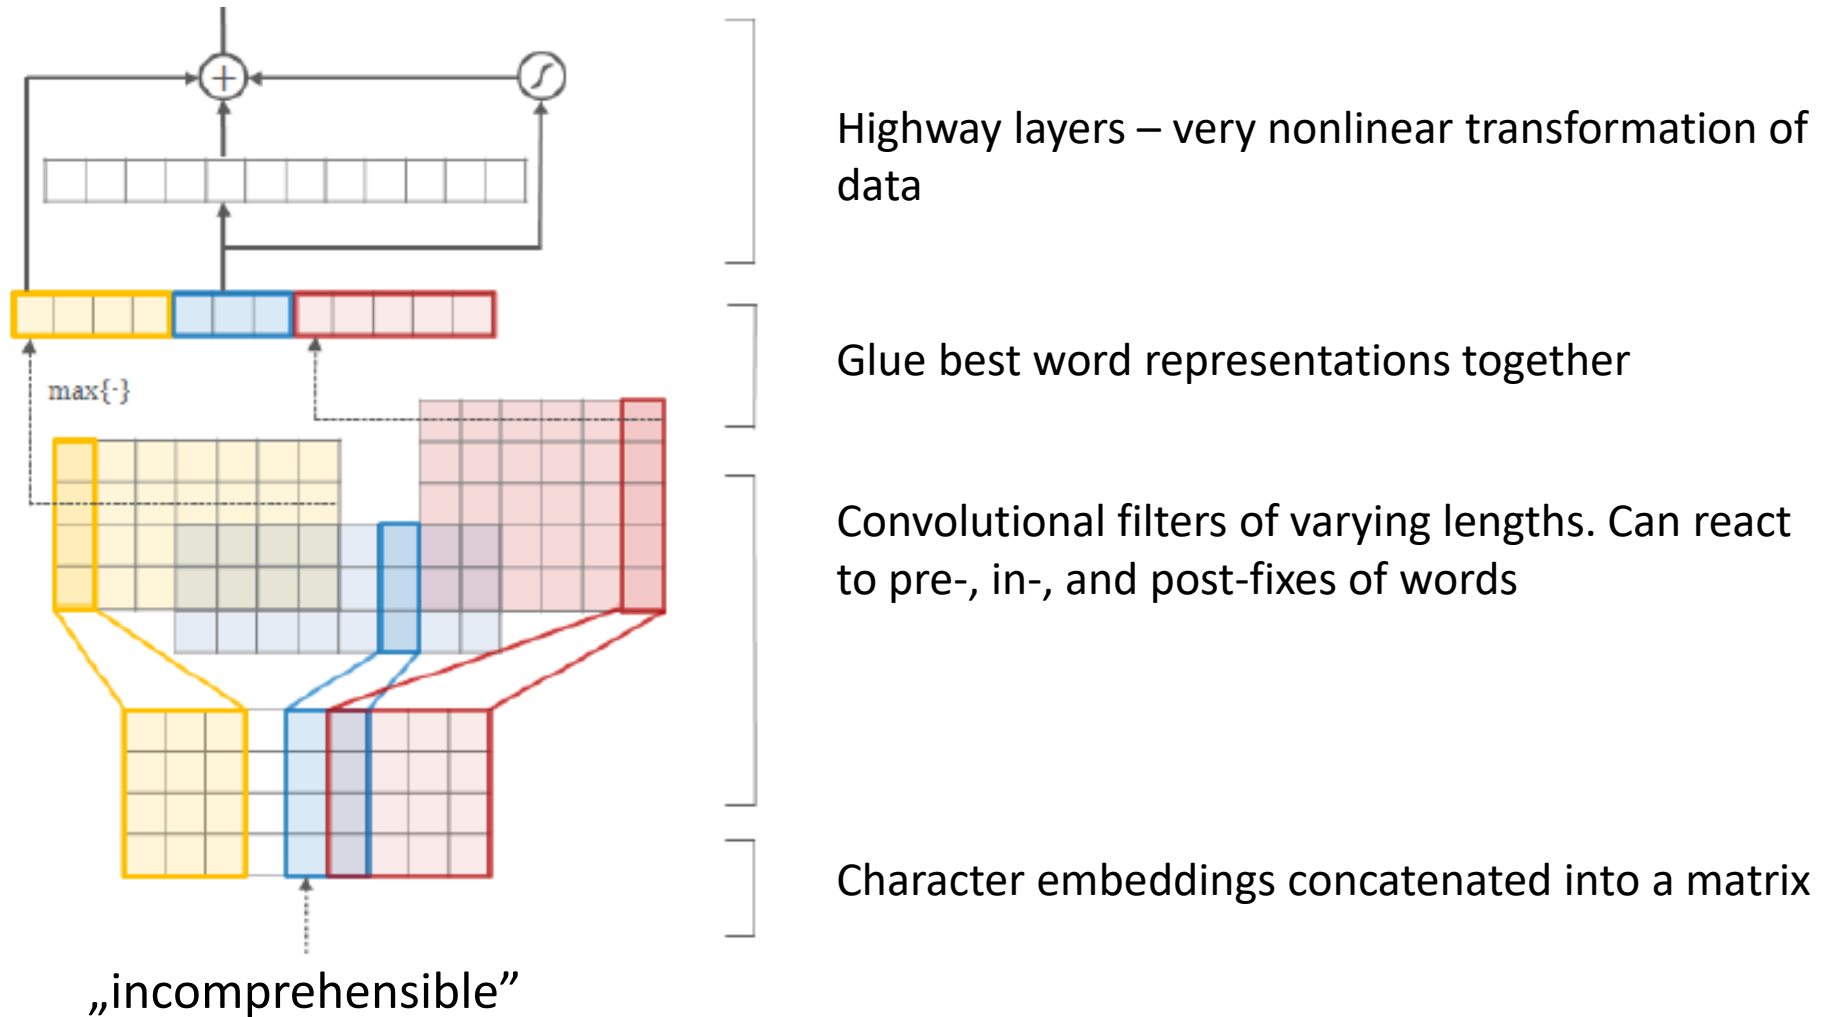
# Dependency parsing



- For each word  $w$
- Two operations:
1. Find head  $h$  (use attention mechanism)
  2. Use  $(w, h)$  to predict dependency type

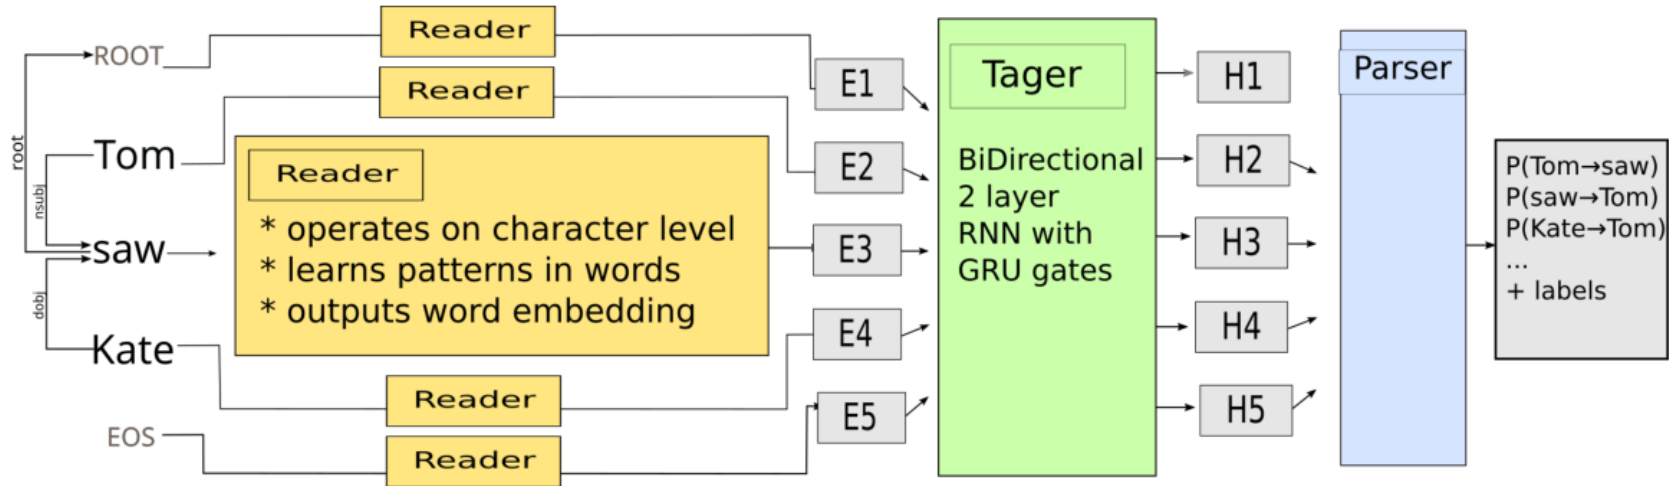


# From characters to word embeddings



Y. Kim, Y. Jernite, D. Sontag, and A. M. Rush, “Character-Aware Neural Language Models,” *arXiv:1508.06615 [cs, stat]*, Aug. 2015.

# From characters to parse trees



**Reader** reads orthographic representations of words and is sensitive to morphemes.

**Tagger** puts words into context

**Parser** finds the dependency edges.

# Jabberwocky (Lewis Carroll)

Tw'as brillig and the slithy toves  
Did gyre and gimble in the wabe;  
All mimsy were the borogoves,  
And the mome raths outgrabe.

# Żabrołak (Stanisław Barańczak)

Brzdęśniało już ślimonne prztowie  
praet:sg:n:perf qub adj:sg:nom:n:pos subst:sg:nom:n

Wyrło i warło się w gulbieży  
praet:sg:n:perf conj praet:sg:n:imperf qub prep:acc:nwok subst:pl:acc:m3

Zmimszałe ćwiły borogowie  
adj:pl:acc:m3:pos praet:pl:f:imperf subst:pl:nom:m1

I rcie grdypały z mrzerzy  
conj subst:pl:nom:n praet:pl:f:imperf prep:gen:nwok subst:sg:gen:f

Underlined words are neologisms, green are correct!

# Multilingual Grammatical Relations

Polish word	Closest russian embeddings
przedwrześniowej	адренергической тренерской таврической непосредственной археологической философской <i>верхнюю</i>
większych	автомобильных <i>трёхдневные</i> технических практических официальных оригинальных
policyjnym	главным историческим глазным непосредственным <i>косыми</i> летним двухсимвольным

- **Green Russian** words have similar grammatical function to **Polish words**.
- **-ской** (skoy) and **-нной** (nnoy) quite distant from polish **-owej** (ovey).
- 3-letter **-ych** paired with 2 letter **-ых**